

Fraud Transaction Detection System - Technical Specifications

1. Project Overview

1.1 Objective

Develop a real-time fraud detection system that classifies transactions as fraudulent or legitimate using machine learning techniques, with the ability to detect multiple fraud scenarios including amount-based, terminal-based, and customer behavior-based fraud patterns.

1.2 Business Value

- **Risk Mitigation:** Prevent financial losses from fraudulent transactions
- **Real-time Protection:** Block suspicious transactions before completion
- **Customer Trust:** Maintain customer confidence through secure transactions
- **Regulatory Compliance:** Meet financial industry fraud prevention requirements
- **Cost Reduction:** Minimize manual review processes and false positives

1.3 Fraud Scenarios to Detect

1. **Amount-based Fraud:** Transactions > \$220 (baseline pattern)
2. **Terminal Compromise:** Fraudulent activity on compromised terminals (28-day windows)
3. **Card-not-Present Fraud:** Customer credential theft with inflated transaction amounts

2. System Architecture

2.1 High-Level Architecture

Transaction Stream → Feature Engineering → ML Model → Risk Scoring → Decision Engine → Action

2.2 System Components

- **Data Ingestion Pipeline:** Real-time transaction processing
- **Feature Engineering Engine:** Creates behavioral and temporal features
- **Model Inference Service:** Fraud prediction and scoring
- **Risk Assessment Module:** Combines model outputs with business rules
- **Alert Management System:** Handles fraud notifications and case management
- **Model Management Platform:** Handles model updates and A/B testing

- **Monitoring Dashboard:** Real-time performance and fraud metrics

3. Data Specifications

3.1 Input Dataset Schema

Feature	Type	Description	Example Values
TRANSACTION_ID	String	Unique transaction identifier	"TXN_001234567"
TX_DATETIME	Datetime	Transaction timestamp	"2024-01-15 14:30:25"
CUSTOMER_ID	String	Customer identifier	"CUST_5678"
TERMINAL_ID	String	Terminal/merchant identifier	"TERM_9012"
TX_AMOUNT	Float	Transaction amount (\$)	125.50
TX_FRAUD	Binary	Fraud label (target)	0 (legitimate), 1 (fraud)

3.2 Data Quality Requirements

- **Completeness:** No missing values in core fields
- **Timeliness:** Transaction data available within 100ms of occurrence
- **Accuracy:** Timestamps must be precise to the second
- **Consistency:** Customer/Terminal IDs follow consistent format

4. Feature Engineering Specifications

4.1 Temporal Features

- **Time-based:**
 - Hour of day, Day of week, Month
 - Is weekend/holiday transaction
 - Time since last transaction (customer/terminal)

4.2 Customer Behavioral Features

- **Spending Patterns** (rolling windows: 1, 7, 14, 30 days):
 - Average transaction amount
 - Standard deviation of amounts
 - Transaction frequency
 - Maximum transaction amount
 - Spending velocity changes

- **Customer Risk Indicators:**
 - Days since first transaction
 - Number of unique terminals used
 - Geographic diversity (if location data available)
 - Ratio of current amount to historical average

4.3 Terminal-based Features

- **Terminal Activity** (rolling windows: 1, 7, 28 days):
 - Transaction volume
 - Average transaction amount
 - Number of unique customers
 - Fraud rate (for model updates)
- **Terminal Risk Indicators:**
 - Terminal age (days since first transaction)
 - Anomalous activity patterns
 - Customer concentration ratio

4.4 Transaction-specific Features

- **Amount Features:**
 - Raw transaction amount
 - Amount rounded to nearest 10/100
 - Log-transformed amount
 - Amount percentile (customer/terminal/global)
- **Derived Features:**
 - Amount deviation from customer average
 - Amount deviation from terminal average
 - Transaction velocity (amount/time since last)

4.5 Aggregated Risk Features

- **Customer Risk Score:** Historical fraud indicators
- **Terminal Risk Score:** Terminal-based risk factors
- **Network Features:** Customer-terminal interaction patterns

5. Machine Learning Model Specifications

5.1 Model Architecture Options

5.1.1 Ensemble Approach (Recommended)

- **Level 1 Models:**
 - Gradient Boosting (XGBoost/LightGBM) - Primary
 - Random Forest - Secondary
 - Logistic Regression - Baseline
- **Level 2 Meta-learner:** Logistic Regression for final prediction

5.1.2 Alternative Single Models

- **Deep Learning:** Neural Network with embedding layers
- **Isolation Forest:** For anomaly detection
- **One-Class SVM:** For outlier detection

5.2 Model Evaluation Metrics

5.2.1 Primary Metrics

- **Precision:** Minimize false positives (legitimate transactions blocked)
- **Recall:** Minimize false negatives (fraudulent transactions missed)
- **F1-Score:** Balance between precision and recall
- **AUC-ROC:** Overall model discrimination ability

5.2.2 Business Metrics

- **False Positive Rate:** < 1% (customer experience)
- **True Positive Rate:** > 95% (fraud detection)
- **Cost-based Evaluation:** Weighted by fraud losses vs. investigation costs

5.3 Model Training Strategy

- **Temporal Split:** Train on historical data, validate on recent data
- **Cross-Validation:** Time-series split respecting temporal order
- **Class Imbalance:** Handle using SMOTE, class weights, or threshold optimization
- **Feature Selection:** Recursive feature elimination and importance analysis

6. Real-time Processing Requirements

6.1 Latency Requirements

- **Model Inference:** < 50ms
- **Feature Engineering:** < 30ms
- **Total Processing Time:** < 100ms per transaction

6.2 Throughput Requirements

- **Peak Load:** 10,000 transactions per second
- **Average Load:** 1,000 transactions per second
- **Scalability:** Auto-scaling based on transaction volume

6.3 Data Pipeline Architecture

```
python
```

```
# Streaming Pipeline Structure
```

```
Transaction Event → Feature Store → Model Inference → Risk Decision → Response
```

7. Technical Implementation

7.1 Technology Stack

- **Programming Language:** Python 3.9+
- **ML Framework:** scikit-learn, XGBoost, LightGBM
- **Stream Processing:** Apache Kafka, Apache Flink
- **Feature Store:** Redis, Apache Cassandra
- **Model Serving:** MLflow, TensorFlow Serving
- **API Framework:** FastAPI
- **Monitoring:** Prometheus, Grafana
- **Containerization:** Docker, Kubernetes

7.2 Code Structure

```
fraud_detection_system/
├── data/
│   ├── raw/
│   ├── processed/
│   └── features/
├── src/
│   ├── data_pipeline/
│   │   ├── ingestion.py
│   │   ├── preprocessing.py
│   │   └── feature_engineering.py
│   ├── models/
│   │   ├── train_model.py
│   │   ├── ensemble_model.py
│   │   └── model_evaluation.py
│   ├── api/
│   │   ├── fraud_api.py
│   │   └── model_serving.py
│   ├── monitoring/
│   │   └── model_monitoring.py
│   └── utils/
│       └── data_utils.py
├── tests/
├── config/
├── models/
├── requirements.txt
└── README.md
```

8. API Specifications

8.1 Fraud Detection Endpoint

python

Real-time Fraud Detection API

POST /api/v1/fraud/detect

```
{
  "transaction_id": "TXN_001234567",
  "customer_id": "CUST_5678",
  "terminal_id": "TERM_9012",
  "tx_amount": 125.50,
  "tx_datetime": "2025-06-23T14:30:25Z"
}
```

Response:

```
{
  "transaction_id": "TXN_001234567",
  "fraud_probability": 0.23,
  "risk_score": 0.15,
  "decision": "APPROVE",
  "risk_factors": [
    "amount_deviation: 0.12",
    "customer_velocity: 0.08"
  ],
  "processing_time_ms": 45,
  "model_version": "v2.1.0",
  "timestamp": "2025-06-23T14:30:25.123Z"
}
```

8.2 Decision Thresholds

- **High Risk:** Probability > 0.8 → BLOCK
- **Medium Risk:** 0.3 < Probability ≤ 0.8 → REVIEW
- **Low Risk:** Probability ≤ 0.3 → APPROVE

9. Model Monitoring and Maintenance

9.1 Performance Monitoring

- **Model Drift Detection:** Feature distribution changes
- **Performance Degradation:** Track precision/recall over time
- **Data Quality Monitoring:** Missing values, outliers, schema changes
- **Latency Monitoring:** Response time tracking

9.2 Retraining Strategy

- **Scheduled Retraining:** Weekly model updates
- **Trigger-based Retraining:** Performance drop > 5%
- **Incremental Learning:** Online learning for rapid adaptation
- **A/B Testing:** Gradual rollout of new models

9.3 Alerting System

- **Critical Alerts:** Model failures, API downtime
- **Performance Alerts:** Degraded accuracy, increased latency
- **Business Alerts:** Fraud spike detection, unusual patterns

10. Security and Compliance

10.1 Data Security

- **Encryption:** Data at rest and in transit
- **Access Control:** Role-based permissions
- **Audit Logging:** Complete transaction trail
- **Data Retention:** Configurable retention policies

10.2 Regulatory Compliance

- **PCI DSS:** Payment card industry standards
- **GDPR:** Data privacy regulations
- **Fair Credit Reporting:** Model explainability
- **Anti-Money Laundering:** Transaction monitoring

11. Testing Strategy

11.1 Model Testing

- **Unit Tests:** Feature engineering functions
- **Integration Tests:** End-to-end pipeline testing
- **Performance Tests:** Load testing, latency validation
- **Shadow Testing:** Compare with existing system

11.2 Fraud Scenario Testing

- **Amount-based:** Transactions > \$220 detection
- **Terminal Compromise:** Multi-day fraud pattern detection
- **Customer Behavior:** Spending anomaly detection
- **Edge Cases:** Boundary condition testing

12. Deployment Strategy

12.1 Infrastructure Requirements

- **Compute:** 8 CPU cores, 32GB RAM per instance
- **Storage:** 1TB SSD for feature store and models
- **Network:** High-bandwidth, low-latency connections
- **Redundancy:** Multi-region deployment for HA

12.2 Deployment Pipeline

Development → Staging → Canary → Production

- **Blue-Green Deployment:** Zero-downtime updates
- **Feature Flags:** Gradual feature rollout
- **Rollback Strategy:** Immediate rollback capability

13. Success Criteria

13.1 Technical KPIs

- **Model Performance:** F1-Score > 0.90
- **System Latency:** < 100ms 99th percentile
- **Availability:** 99.99% uptime
- **Throughput:** Handle peak loads without degradation

13.2 Business KPIs

- **Fraud Detection Rate:** > 95%
- **False Positive Rate:** < 1%
- **Cost Savings:** Measurable reduction in fraud losses
- **Customer Satisfaction:** Minimal impact on legitimate transactions

14. Risk Management

14.1 Technical Risks

- **Model Bias:** Regular bias testing and mitigation
- **Adversarial Attacks:** Fraud pattern evolution
- **System Failures:** Fallback to rule-based systems
- **Data Quality Issues:** Robust data validation

14.2 Business Risks

- **Regulatory Changes:** Compliance monitoring
- **Competitive Pressure:** Continuous improvement
- **Fraud Evolution:** Adaptive learning capabilities
- **Customer Impact:** Careful threshold management

15. Future Enhancements

15.1 Advanced Features

- **Graph Neural Networks:** Customer-merchant relationship analysis
- **Reinforcement Learning:** Adaptive decision making
- **Federated Learning:** Privacy-preserving model updates
- **Real-time Feature Engineering:** Stream processing optimization

15.2 Integration Opportunities

- **External Data Sources:** Credit bureaus, fraud databases
- **Multi-channel Analysis:** Cross-channel fraud patterns
- **Behavioral Biometrics:** User interaction patterns
- **Geographic Intelligence:** Location-based risk assessment